

CSC 480/580 Principles of Machine Learning

01 Supervised learning; Decision Trees



Jason Pacheco

Department of Computer Science



The supervised learning problem

Supervised learning setup: putting it together



- Goal: design learning algorithm $\mathcal{A}$, such that:
on iid training data S , its output f has low *generalization error* $L_D(f)$


Quiz 1

Which of the following statements are part of the standard supervised learning protocol?

- A. The training and test datasets are drawn from the same underlying data distribution. Yes
- B. The training and test datasets do not contain any overlapping data points. Yes
- C. The loss function is chosen to reflect the prediction task and how errors should be measured. Yes
- D. The test dataset is used during training to improve the learned model. No
- E. The loss function can be chosen independently of the task without affecting model evaluation. No

First supervised learning algorithm:
learning decision trees

Supervised Learning



Goal Learn function f from training data that makes predictions on unseen test data

Prototypical supervised learning problems:

- Regression
- Classification (binary / multiclass)
- Ranking
- ...

This lecture will focus on
binary classification...

Example: course recommendation

Task *Given a student, recommend a set of courses that s/he would like*

We are allowed to ask a sequence of questions...

You: Is the course under consideration in Systems?

Me: Yes

You: Has this student taken any other Systems courses?

Me: Yes

You: Has this student liked most previous Systems courses?


Me: No

You: *I predict this student will not like this course.*

The Machine Learning approach:

learned from previous students' course survey results

course description
student info.



Model: Decision Tree

Use our questions to build a binary tree:

You: Is the course under consideration in Systems?

Me: Yes

You: Has this student taken any other Systems courses?

Me: Yes

You: Has this student liked most previous Systems courses?

Me: No

You: *I predict this student will not like this course.*

Terminology:

- Question & Answer → Feature
- Set of Question & Answers → Training Data
- “Like” / “Not-like” → Labels




Figure 1.2: A decision tree for a course recommender system, from which the in-text “dialog” is drawn.

Training Dataset: previous course survey results

Define the labeled training dataset $S = \{(x_i, y_i)\}_{i=1}^m$

To make this a binary classification we set
“Like” = $\{+2, +1, 0\}$
“Not-like” = $\{-1, -2\}$

	Rating	Easy?	AI?	Sys?	Thy?	Morning?
	+2	y	y	n	y	n
	+2	y	y	n	y	n
	+2	n	y	n	n	n
	+2	n	n	n	y	n
	+2	n	y	y	n	y
	+1	y	y	n	n	n
	+1	y	y	n	y	n
	+1	n	y	n	y	n
	0	n	n	n	n	y
	0	y	n	n	y	y
	0	n	y	n	y	n
	0	y	y	y	y	y
	-1	y	y	y	n	y
	-1	n	n	y	y	n
	-1	n	n	y	n	y
	-1	y	n	y	n	y
	-2	n	n	y	y	n
	-2	n	y	y	n	y
	-2	y	n	y	n	n
	-2	y	n	y	n	y

Decision trees: basic terminology




Figure 1.2: A decision tree for a course recommender system, from which the in-text “dialog” is drawn.

node
root node
leaf node
internal node

parent
children
subtree
depth

- Key advantage of using decision trees for decision making: *intepretability*
- Useful in consequential settings, e.g. medical treatment, loan approval, etc.


nodes organized in a tree-based structure, leading to a prediction (Fig. 1). The interpretability of decision trees allows physicians to understand why a prediction or stratification is being made, providing an account of the reasons behind the decision to subsequently accept or override the model's output. This interaction between humans and algorithms can provide

Prediction using decision trees

- Test: predict using a decision tree:

Algorithm 2 `DECISIONTREETEST`(*tree*, *test point*)

```
1: if tree is of the form LEAF(guess) then  
2:   return guess  
3: else if tree is of the form NODE(f, left, right) then  
4:   if f = no in test point then  
5:     return DECISIONTREETEST(left, test point)  
6:   else  
7:     return DECISIONTREETEST(right, test point)  
8:   end if  
9: end if
```



guess=prediction

left=no
right=yes

- Training: apply a learning algorithm that can build trees f from training data
- How to design the algorithm?

Learning Decision Trees

Example Guess a number between 1 and 100. Which set of questions are better? Why?

Set 1

- 1) Greater than 20? **Y**
- 2) Has a 7 in it? **N**
- 3) Odd? **N**

Set 2

- 1) Greater than 50? **Y**
- 2) Greater than 75? **N**
- 3) Greater than 63? **N**

How many questions should this problem require?

Key Intuition: The decision tree should try to ask informative questions


Training accuracy / error

- The training data $S = \{(x_i, y_i)\}_{i=1}^m$
- Predictor f 's training accuracy = fraction of examples in S that are **correctly classified** by f
- In formula,

$$A_S(f) = \frac{1}{m} \sum_{i=1}^m I(f(x_i) = y_i) \quad I(A) = 1 \text{ if } A \text{ happens; } =0 \text{ otherwise}$$

- Training error $L_S(f) = 1 - A_S(f) = \frac{1}{m} \sum_{i=1}^m I(f(x_i) \neq y_i)$
- The training error minimization approach:
 - Idea: f has low $L_S(f) \Rightarrow f$ likely has low $L_D(f)$
 - Approach: Find f with lowest $L_S(f) \Leftrightarrow f$ with highest $A_S(f)$
 - Problem with this approach?

training accuracy = 47 / 50



May overfit – we'll come back to this

Decision tree training: zero-level case

- Q: if I would like to find a depth-0 tree with the highest training accuracy, what would that tree look like?
- The tree will not have internal nodes (no questions asked), and just has a leaf





- Choosing leaf to be "like" – training accuracy = 60%
- Choosing leaf to be "not-like" – training accuracy = 40%
- The optimal leaf would be "like"

Decision tree training: single level case

- Q: if I could only ask one question (design a depth-1 tree), what question would I ask?
- Perhaps useful: look at the *histograms* of labels (for **like** & **not-like**), grouped by values taken by the feature




- Example: feature “easy” vs. feature “AI”




Histogram calculation

Rating	Easy?	AI?	Sys?	Thy?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
o	n	n	n	n	y
o	y	n	n	y	y
o	n	y	n	y	n
o	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y



Decision tree training: single level case


- Q: if I could only ask one question (design a depth-1 tree), what question would I ask?
- Intuition: look at the histograms of labels for each feature
- Which feature (question) is better, 'easy' or 'AI'?
- Best training accuracy using 'easy':
 $(\max(6,4) + \max(6,4)) / 20 = 0.6$
- Best training accuracy using 'AI':
 $(\max(9,2) + \max(3,6)) / 20 = 0.75$
- Thus, 'AI' is more informative (more useful in telling apart like & not-like)



Decision tree training: single level case

- In formula:

$$\begin{aligned} \text{Score}(f, S) := & \max\left(P_S(y = +, x_f = \text{yes}), P_S(y = -, x_f = \text{yes})\right) \\ & + \max\left(P_S(y = +, x_f = \text{no}), P_S(y = -, x_f = \text{no})\right) \end{aligned}$$



- Here, $P_S(E)$ denotes the probability of event E if an example (x, y) is chosen uniformly from S
- In other words: the frequency of E in sample S
- E.g.

$$\text{Score}(\text{'easy'}, S) = \max\left(\frac{6}{20}, \frac{4}{20}\right) + \max\left(\frac{6}{20}, \frac{4}{20}\right) = 0.6$$

Math interlude: Uniform distribution over a set

Given set $S = \{v_1, v_2, \dots, v_N\}$; Z is drawn from the uniform distribution of S , then

$$P(Z = k) = \begin{cases} \frac{1}{N}, & k \in \{v_1, v_2, \dots, v_N\} \\ 0, & \text{otherwise} \end{cases}$$

We write this as $Z \sim \text{Uniform}(S)$

- Selecting a student from a class
- Choosing an example from a dataset
- $P_{Z \sim \text{Uniform}(S)}$ is often abbreviated as P_S
- $P_S(E)$ = fractions of elements in S that satisfies E



airplane



automobile




bird



Decision tree training: single level case

- In formula:

$$\text{Score}(f, S) := \max(P_S(y = +, x_f = \text{yes}), P_S(y = -, x_f = \text{yes})) \\ + \max(P_S(y = +, x_f = \text{no}), P_S(y = -, x_f = \text{no}))$$



- Written in conditional probability form: A measure of label “purity”

$$= \max(P_S(y = + | x_f = \text{yes}), P_S(y = - | x_f = \text{yes})) \cdot P_S(x_f = \text{yes}) \\ + \max(P_S(y = + | x_f = \text{no}), P_S(y = - | x_f = \text{no})) \cdot P_S(x_f = \text{no})$$





$$P(AB) = P(A)P(B | A)$$

- e.g. $\text{Score}(\text{'easy'}, S) = \max(0.6, 0.4) \times 0.5 + \max(0.6, 0.4) \times 0.5 = 0.6$

- Interpretation: $\text{Score}(f, S)$ measures the ability of feature f in predicting label y – **informativeness** of feature f

Decision tree training: general level case

- Key idea: divide & conquer
- Build the root of the tree greedily
- Build the left and right subtrees *recursively*
- When to stop the recursion?




Algorithm 1 `DECISIONTREETRAIN`(*data*, *remaining features*)

```
1: guess ← most frequent answer in data           // default answer for this data
2: if the labels in data are unambiguous then
3:   return LEAF(guess)                          // base case: no need to split further
4: else if remaining features is empty then
5:   return LEAF(guess)                          // base case: cannot split further
6: else                                             // we need to query more features
7:   for all  $f \in \textit{remaining features}$  do
8:     NO ← the subset of data on which  $f=no$ 
9:     YES ← the subset of data on which  $f=yes$ 
10:    score[f] ← # of majority vote answers in NO
11:                + # of majority vote answers in YES
12:                // the accuracy we would get if we only queried on f
13:   end for
14:   f ← the feature with maximal score(f)
15:   NO ← the subset of data on which  $f=no$ 
16:   YES ← the subset of data on which  $f=yes$ 
17:   left ← DECISIONTREETRAIN(NO, remaining features \ {f})
18:   right ← DECISIONTREETRAIN(YES, remaining features \ {f})
19:   return NODE(f, left, right)
20: end if
```

answer=label

Unambiguous = all identical



$\text{Score}(f, S)$ = informativeness of feature f for dataset S

Q: does this algorithm always terminate?

A: Yes, since each recursive call decreases the number of available features by 1

Dealing with various types of features

- Binary: $x_f \in \{0,1\}$

- Node: $x_f = 0$?

Easy?	AI?	Sys?	Thy?	Morning?
y	y	n	y	n


- Categorical: $x_f \in \{1,2, \dots, C\}$

- Node: $x_f \in \{i_1, \dots, i_l\}$?



- real value: $x_f \in \mathbb{R}$

- Node: $x_f \leq z$?



Example: spam filtering I

- ▶ Spam dataset
- ▶ 4601 email messages, about 39% are spam
- ▶ Classify message by spam and not-spam
- ▶ 57 features
 - ▶ 48 are of the form “percentage of email words that is (WORD)”
 - ▶ 6 are of the form “percentage of email characters is (CHAR)”
 - ▶ 3 other features (e.g., “longest sequence of all-caps”)
- ▶ Final tree after pruning has 17 leaves, 9.3% test error rate


The meaning of numbers: e.g., 600 / 1536

test examples passing
this node with minority label



test examples passing
through this node

Q: What is the best depth-0 decision tree, and what is its accuracy?

A: predicting “email” for all examples, accuracy = $736 / 1536 = 47.9\%$



Decision tree training: top-down approach




Algorithm 1 `DECISIONTREETRAIN`(*data*, *remaining features*)

```
1: guess  $\leftarrow$  most frequent answer in data // default answer for this data
2: if the labels in data are unambiguous then
3:   return LEAF(guess) // base case: no need to split further
4: else if remaining features is empty then
5:   return LEAF(guess) // base case: cannot split further
6: else // we need to query more features
7:   for all  $f \in \text{remaining features}$  do
8:     NO  $\leftarrow$  the subset of data on which  $f=no$ 
9:     YES  $\leftarrow$  the subset of data on which  $f=yes$ 
10:    score[f]  $\leftarrow$  # of majority vote answers in NO
11:    + # of majority vote answers in YES
12:    // the accuracy we would get if we only queried on f
13:   end for
14:   f  $\leftarrow$  the feature with maximal score(f)
15:   NO  $\leftarrow$  the subset of data on which  $f=no$ 
16:   YES  $\leftarrow$  the subset of data on which  $f=yes$ 
17:   left  $\leftarrow$  DECISIONTREETRAIN(NO, remaining features  $\setminus$  {f})
18:   right  $\leftarrow$  DECISIONTREETRAIN(YES, remaining features  $\setminus$  {f})
19:   return NODE(f, left, right)
20: end if
```

answer=label

Unambiguous = all identical




$\text{Score}(f, S)$ = informativeness of feature f for dataset S

Q: does this algorithm always terminate?

A: Yes, since each recursive call decreases the number of available features by 1

Example: spam filtering I

- ▶ Spam dataset
- ▶ 4601 email messages, about 39% are spam
- ▶ Classify message by spam and not-spam
- ▶ 57 features
 - ▶ 48 are of the form “percentage of email words that is (WORD)”
 - ▶ 6 are of the form “percentage of email characters is (CHAR)”
 - ▶ 3 other features (e.g., “longest sequence of all-caps”)
- ▶ Final tree after pruning has 17 leaves, 9.3% test error rate



The meaning of numbers: e.g., 600 / 1536

test examples passing
this node with minority label

test examples passing
through this node



Q: What is the best depth-0 decision tree, and what is its accuracy?

A: predicting “email” for all examples, accuracy = $736 / 1536 = 47.9\%$

Decision tree training: generalized informativeness scores

- $\text{Score}(f)$ = a measure of **informativeness** of f
- Alternative view -- **uncertainty reduction**: how much uncertainty about y can be reduced if we know f
- Uncertainty measures of population:

We used classif. accuracy:
 $\max(p, 1-p)$,



Notions of uncertainty: binary case ($\mathcal{Y} = \{0, 1\}$)

Suppose in a set of examples $S \subseteq \mathcal{X} \times \{0, 1\}$, a p fraction are labeled as 1

① **Classification error:**


$$u(S) := \min\{p, 1 - p\}$$

② **Gini index:**

$$u(S) := 2p(1 - p)$$

③ **Entropy:**

$$u(S) := p \log \frac{1}{p} + (1-p) \log \frac{1}{1-p}$$




(3) is divided by 2
so the plot looks
comparable

log here is base-2

Decision tree training: generalized informativeness scores

Suppose the data S at a leaf ℓ is split by a rule h into S_L and S_R , where $p_L := |S_L|/|S|$ and $p_R := |S_R|/|S|$





The **reduction in uncertainty** from using rule h at leaf ℓ is

$$u(S) - \left(p_L \cdot u(S_L) + p_R \cdot u(S_R) \right) \boxed{=: \text{Score}(h, S) \text{ (Generalized)}}$$

Decision tree training: generalized informativeness scores

- Multiclass classification setting: $\mathcal{Y} = \{1, \dots, K\}$



Notions of uncertainty: general case

Suppose in $S \subseteq \mathcal{X} \times \mathcal{Y}$, a p_k fraction are labeled as k (for each $k \in \mathcal{Y}$).

- 1 **Classification error:**

$$u(S) := 1 - \max_{k \in \mathcal{Y}} p_k$$

- 2 **Gini index:**

$$u(S) := 1 - \sum_{k \in \mathcal{Y}} p_k^2$$

- 3 **Entropy:**

$$u(S) := \sum_{k \in \mathcal{Y}} p_k \log \frac{1}{p_k}$$


Each is *maximized* when $p_k = 1/|\mathcal{Y}|$ for all $k \in \mathcal{Y}$
(i.e., equal numbers of each label in S)

Each is *minimized* when $p_k = 1$ for a single label $k \in \mathcal{Y}$
(so S is **pure** in label)

Entropy Uncertainty

Entropy of random variable Y : $H(Y) = \sum_y P(Y = y) \log \frac{1}{P(Y=y)}$

Coin Flip Example: $Y \sim \text{Bernoulli}(p)$



- Key object studied in information & coding theory
- Interpretations:
 - the minimum number of bits needed to reliably encode Y
 - the uncertainty about outcome of Y

Maximum uncertainty when coin is fair.

Information Theory

- Shannon's seminal paper, "A Mathematical Theory of Communication", Bell Labs Technical Journal, 1948
- Claude Shannon (1916-2001): pioneered AI research

corridors until it found the target.^[65] Having travelled through the maze, the mouse could then be placed anywhere it had been before, and because of its prior experience it could go directly to the target. If placed in unfamiliar territory, it was programmed to search until it reached a known location and then it would proceed to the target, adding the new knowledge to its memory and learning new behavior.^[65] After much trial and error, this device would learn the shortest
- "Prediction and Entropy of Printed English" – early work on the n-gram model in NLP



Generalized informativeness score in action


- Example: consider entropy uncertainty, on dataset S

$$u(S) = \sum_{y_0 \in \{+, -\}} P_S(y = l) \log \frac{1}{P_S(y=l)}$$

- Score(S, f) = $u(S) - (p_L u(S_L) + p_R u(S_R))$

$$P_S(x_f = no)$$

$$\sum_{l \in \{+, -\}} P_S(y = l \mid x_f = no) \log \frac{1}{P_S(y = l \mid x_f = no)}$$



- In the above example:

- $u(S) = 0.6 \log \frac{1}{0.6} + 0.4 \log \frac{1}{0.4},$
- $p_L = 0.5, p_R = 0.5$
- $u(S_L) = u(S_R) = 0.6 \log \frac{1}{0.6} + 0.4 \log \frac{1}{0.4}$
 $\Rightarrow \text{Score}(' \text{ easy}', S) = 0$

- In this case, Score(S, f) is also known as the mutual information between x_f and y (under P_S)

Mutual Information

$$I(X; Y) = H(X) - H(X | Y) \quad \text{Recall: } H(X) = \sum_x P(X = x) \log \frac{1}{P(X=x)}$$


- Measures entropy reduction after observing Y
- How much information does Y carry about X?

$$\text{Aside: } H(X|Y) = \sum_{x,y} P(X = x, Y = y) \log \frac{1}{P(X=x|Y=y)}$$

Stop splitting when there is no reduction in uncertainty? This is a bad idea!

The 'XOR' data: Suppose $\mathcal{X} = \mathbb{R}^2$ and $\mathcal{Y} = \{\text{red, blue}\}$, and the data is as follows:


A	B	A XOR B
0	0	0
0	1	1
1	0	1
1	1	0




- Any axis-aligned split has no reduction on uncertainty on S
- However, a depth-2 decision tree (with axis-aligned splits) has zero training error

Overfitting can happen

- “Spurious” patterns can be learned.



A better alternative:



- A fix to the training algorithm: stop recursion & return leaf once we reach depth k (say $k = 2$)
- Alternatively: do **pruning** on trained decision tree – a popular heuristic

Next lecture

- Supervised learning: what to do if the data distribution is *known*?
- Models, parameters, hyperparameters
- Practical considerations
- Assigned reading: CIML Chap. 2 (Limits of learning)